

```
In [ ]:

In [ ]: # Build an end-to-end prediction project: Artificial Neural Networks (ANNs) for Natural Language Processing (NLP) with predication
# Showing my ANNs skills

In [ ]: # Import Libraries
import pickle
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder

# ANN
import datetime as t
import tensorflow as tf
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Sequential
from tensorflow.keras.callbacks import TensorBoard, EarlyStopping
from tensorflow.keras.callbacks import EarlyStopping, TensorBoard
from tensorflow.keras.models import load_model

In [ ]: # read "csv" file
datac=pd.read_csv("ann_project.csv")

In [ ]: # look at first 5 rows in data
datac.head()

In [ ]: # Dataset processing

# Drop some columns ([ 'col_1', 'col_2', 'surname'],axis=1) which do not play key roles for predication.

datach=datac.drop([ 'col_1', 'col_2', 'surname'],axis=1)
datach

In [ ]: # Use LabelEncoder() for "'Gender' column
LE_Gender=LabelEncoder()
datach['Gender']=LE_Gender.fit_transform(datach['Gender'])
datach

In [ ]: # save results from =LabelEncoder() used on 'Gender' column
with open('LE_Gender.pkl','wb') as gen:
    pickle.dump(LE_Gender,gen)

In [ ]: # use OneHotEncoder () for 'Geography' column
OHEG=OneHotEncoder()
oheg=OHEG.fit_transform(datach[['Geography']]).toarray()
oheg

In [ ]: # save results from OneHotEncoder used on 'Geography' column
with open('OHEG.pkl','wb') as geo:
    pickle.dump(OHEG,geo)

In [ ]: # Find new feature names in 'Geography' column
OHEG.get_feature_names_out(['Geography'])

In [ ]: # use OneHotEncoder () features to create a DataFrame
geography_dataframe = pd.DataFrame(oheg,columns=OHEG.get_feature_names_out(['Geography']))
geography_dataframe

In [ ]: # deop 'Geography' column from the original dataset, and then add geography_dataframe
data_concat=pd.concat([datach.drop('Geography',axis=1),geography_dataframe,axis=1)

In [ ]: data_concat.head()

In [ ]: # use the "data_concat" dataset to find the independent (y) columns and dependent (x) columns.

x=data_concat.drop('Exited',axis=1) # Independent features
y=data_concat['Exited'] # Dependent feature (target)

#print(x)
#print(y)

In [ ]: # Dive the dataset into in train and test sets
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=123)

In [ ]: # Use StandardScaler() to scale x and y columns.
standard_scaler=StandardScaler()

x_train=standard_scaler.fit_transform(x_train)

x_test=standard_scaler.transform(x_test) # transform() uses results received from the fit_transform() on x_train.

In [ ]: # store the fitted results (mean, and standard deviation) received from x_train datasets.

with open('standard_scaler.pkl','wb') as ss:
    pickle.dump(standard_scaler, ss)

Artificial Neural Network (ANN) model

In [ ]: # Finding number of featurers.
(x_train.shape[1],)

In [ ]: # making an ANN model

model = Sequential([Input(shape=(x_train.shape[1],)), Dense(128, activation='relu'), Dense(64, activation='relu'), Dense(1, activation='sigmoid')])

In [ ]: # model summary
model.summary()

In [ ]: # Use Optimizer to update weights in model.

optimizer_ann=tf.keras.optimizers.Adam(learning_rate=0.01) # use default learning rate
optimizer_ann

In [ ]: # find error
loss_ann=tf.keras.losses.BinaryCrossentropy()
loss_ann

In [ ]: # compile the model with optimize, loss, and accuracy.

model.compile(optimizer=optimizer_ann,loss= loss_ann,metrics=['accuracy'])

In [ ]: # Build the Tensorboard for monitoring pgogress

import datetime

log_dir_fit="logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")

tensorboard_fit=TensorBoard(log_dir=log_dir_fit,histogram_freq=1)

In [ ]: # Set up Early Stopping to stop overfitting
early_stopping_fit=EarlyStopping(monitor='val_loss',patience=10,restore_best_weights=True)

In [ ]: # run the model
train_model=model.fit(x_train,y_train,validation_data=(x_test,y_test),epochs=200, callbacks=[tensorboard_fit, early_stopping_fit])

In [ ]: # Save model
model.save('ann_model.h5')

In [ ]: # For visualization, load tensorboard
!load_ext tensorboard

In [ ]: # Run tensorboard on port (6007)
!tensorboard --logdir logs/fit --port=6008

Prediction

In [ ]: # input example dataset

sample = {
    'CreditScore': 800,
    'Geography': 'Spain',
    'Gender': 'Female',
    'Age': 46,
    'Tenure': 8,
    'Balance': 100000,
    'NumOfProducts': 4,
    'HasCrCard': 0,
    'IsActiveMember': 0,
    'EstimatedSalary': 110000
}

In [ ]: # Load the trained model
model=load_model('ann_model.h5')

In [ ]: # load saved files

# geography
with open('OHEG.pkl','rb') as geo:
    LE_geo=pickle.load(geo)

# gender
with open('LE_Gender.pkl','rb') as gen:
    LE_gen = pickle.load(gen)

# scalar
with open('standard_scaler.pkl','rb') as ss:
    scal = pickle.load(ss)

In [ ]: # One-hot encode 'Geography'
sample_geography = pd.DataFrame({'Geography': [sample['Geography']]})

In [ ]: sample_geography_encoded = OHEG.transform(sample_geography).toarray()

In [ ]: sample_geography_encoded_data = pd.DataFrame(sample_geography_encoded, columns=OHEG.get_feature_names_out(['Geography']))
sample_geography_encoded_data

In [ ]: # Use input (sample) data to create DataFrame
sample_data=pd.DataFrame([sample])
sample_data

In [ ]: # Encode 'Gender' Using LabelEncoder
sample_data['Gender']=LE_Gender.transform(sample_data['Gender'])
sample_data

In [ ]: # concatinatne one hot encoded "Geography" to input DataFrame.

sample_data=pd.concat([sample_data.drop("Geography",axis=1), sample_geography_encoded_data,axis=1)
sample_data

In [ ]: # Scaling (normalizing) the input data

sample_data_scale=standard_scaler.transform(sample_data)
sample_data_scale

In [ ]: # Predict sample
sample_data_prediction=model.predict(sample_data_scale)
sample_data_prediction

In [ ]: # sample prediction probability
if sample_data_prediction > 0.5:
    print('customer is going to leave.')
else:
    print('customer is going to stay.')

In [ ]:

In [ ]: """
conclusion:

Designed and trained a customer churn prediction model using an ANN in TensorFlow/Keras, resulting ~ 87% validation accuracy with early stopping on a dataset. Built an end-to-end prediction system that results ~99% churn probability for a sample case using persisted encoders and scalers with scikit-learn.

"""
```

